

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/327866113>

Prior Exposure Increases Perceived Accuracy of Fake News

Article in *Journal of Experimental Psychology General* · September 2018

DOI: 10.1037/xge0000465

CITATIONS

804

READS

9,702

3 authors, including:



Gordon Pennycook
Cornell University

213 PUBLICATIONS 25,165 CITATIONS

[SEE PROFILE](#)



Tyrone Cannon
Yale University

839 PUBLICATIONS 49,821 CITATIONS

[SEE PROFILE](#)

Journal of Experimental Psychology: General

Prior Exposure Increases Perceived Accuracy of Fake News

Gordon Pennycook, Tyrone D. Cannon, and David G. Rand

Online First Publication, September 24, 2018. <http://dx.doi.org/10.1037/xge0000465>

CITATION

Pennycook, G., Cannon, T. D., & Rand, D. G. (2018, September 24). Prior Exposure Increases Perceived Accuracy of Fake News. *Journal of Experimental Psychology: General*. Advance online publication. <http://dx.doi.org/10.1037/xge0000465>

Prior Exposure Increases Perceived Accuracy of Fake News

Gordon Pennycook, Tyrone D. Cannon, and David G. Rand
Yale University

The 2016 U.S. presidential election brought considerable attention to the phenomenon of “fake news”: entirely fabricated and often partisan content that is presented as factual. Here we demonstrate one mechanism that contributes to the believability of fake news: fluency via prior exposure. Using actual fake-news headlines presented as they were seen on Facebook, we show that even a single exposure increases subsequent perceptions of accuracy, both within the same session and after a week. Moreover, this “illusory truth effect” for fake-news headlines occurs despite a low level of overall believability and even when the stories are labeled as contested by fact checkers or are inconsistent with the reader’s political ideology. These results suggest that social media platforms help to incubate belief in blatantly false news stories and that tagging such stories as disputed is not an effective solution to this problem. It is interesting, however, that we also found that prior exposure does not impact entirely implausible statements (e.g., “The earth is a perfect square”). These observations indicate that although extreme implausibility is a boundary condition of the illusory truth effect, only a small degree of potential plausibility is sufficient for repetition to increase perceived accuracy. As a consequence, the scope and impact of repetition on beliefs is greater than has been previously assumed.

Keywords: fake news, news media, social media, fluency, illusory truth effect

Supplemental materials: <http://dx.doi.org/10.1037/xge0000465.supp>

The ability to form accurate beliefs, particularly about issues of great importance, is key to people’s success as individuals as well as the functioning of their societal institutions (and, in particular, democracy). Across a wide range of domains, it is critically important to correctly assess what is true and what is false: Accordingly, differentiating real from unreal is at the heart of society’s constructs of rationality and sanity (Corlett, Krystal, Taylor, & Fletcher, 2009; Sanford, Veckenstedt, Moritz, Balzan, & Woodward, 2014). Yet the ability to form and update beliefs about the world sometimes goes awry—and not just in the context of inconsequential, small-stakes decisions.

The potential for systematic inaccuracy in important beliefs has been particularly highlighted by the widespread consumption of disinformation during the 2016 U.S. presidential election. This is most notably exemplified by so-called fake news—that is, news stories that were fabricated (but presented as if from legitimate sources) and promoted on social media to deceive the public for ideological and/or financial gain (Lazer et al., 2018). An analysis of the top performing news articles on Facebook in the months leading up to the election revealed that the top fake-news articles actually outperformed the top real-news articles in terms of shares, likes, and comments (Silverman, Strapagiel, Shaban, & Hall,

Gordon Pennycook and Tyrone D. Cannon, Department of Psychology, Yale University; David G. Rand, Department of Psychology, Department of Economics, and School of Management, Yale University.

David G. Rand is now at the Sloan School of Management and Department of Brain & Cognitive Sciences, Massachusetts Institute of Technology.

All data are available online (<https://osf.io/txf46/>). A working paper version of the article was posted online via the Social Sciences Research Network (<https://ssrn.com/abstract=2958246>) and on ResearchGate (https://www.researchgate.net/publication/317069544_Prior_Exposure_Increases_Perceived_Accuracy_of_Fake_News). Portions of this research were presented in 2017 and 2018 at the following venues: (a) Conference for Combating Fake News: An Agenda for Research and Action, Harvard Law School and Northeastern University; (b) Canadian Society for Brain, Behaviour and Cognitive Science Annual Conference, in Regina, Saskatchewan; (c) Brown University, Department of Psychology; (d) Harvard University, Department of Psychology; (e) Yale University, Institute for

Network Science and Technology and Ethics Study Group; (f) University of Connecticut, Cognitive Science Colloquium; (g) University of Regina, Department of Psychology; (h) University of Saskatchewan, Department of Psychology; and (i) the 2018 Annual Convention of the Society of Personality and Social Psychology, in Atlanta, Georgia.

This research was supported by a Social Sciences and Humanities Council of Canada Banting Postdoctoral Fellowship (to Gordon Pennycook), National Institute of Mental Health Grant MH081902 (to Tyrone D. Cannon), and grants from the Templeton World Charity Foundation and the Defense Advanced Research Projects Agency (to David G. Rand). The content is solely the responsibility of the authors and does not necessarily reflect the official views of the Social Sciences and Humanities Council of Canada, the National Institute of Mental Health, or the Defense Advanced Research Projects Agency.

Correspondence concerning this article should be addressed to Gordon Pennycook, who is now at the Faculty of Business Administration, University of Regina, Regina, Saskatchewan, Canada, S4S 0A2. E-mail: grpennycook@gmail.com

2016). Although it is unclear to what extent fake news influenced the outcome of the presidential election (Allcott & Gentzkow, 2017), there is no question that many people were deceived by entirely fabricated (and often quite fanciful) fake-news stories—people including, for example, high-ranking government officials, such as Pakistan’s defense minister (Goldman, 2016). How is it that so many people came to believe stories that were patently and demonstrably untrue? What mechanisms underlie these false beliefs that might be called mass delusions?

Here, we explore one potential answer: prior exposure. Given the ease with which fake news can be created and distributed on social media platforms (Shane, 2017), combined with the increasing tendency to consume news via social media (Gottfried & Shearer, 2016), it is likely that people are being exposed to fake-news stories with much greater frequency than in the past. Might exposure per se help to explain people’s tendency to believe outlandish political disinformation?

The Illusory Truth Effect

There is a long tradition of work in cognitive science demonstrating that prior exposure to a statement (e.g., “The capybara is the largest of the marsupials”) increases the likelihood that participants will judge it to be accurate (Arkes, Boehm, & Xu, 1991; Bacon, 1979; Begg, Anas, & Farinacci, 1992; Dechêne, Stahl, Hansen, & Wänke, 2010; Fazio, Brashier, Payne, & Marsh, 2015; Hasher, Goldstein, & Toppino, 1977; Polage, 2012; Schwartz, 1982). The dominant account of this “illusory truth effect” is that repetition increases the ease with which statements are processed (i.e., processing fluency), which in turn is used heuristically to infer accuracy (Alter & Oppenheimer, 2009; Begg et al., 1992; Reber, Winkielman, & Schwarz, 1998; Unkelbach, 2007; Wang, Brashier, Wing, Marsh, & Cabeza, 2016; Whittlesea, 1993; but see Unkelbach & Rom, 2017). Past studies have shown this phenomenon using a range of innocuous and plausible statements, such as obscure trivia questions (Bacon, 1979) or assertions about consumer products (Hawkins & Hoch, 1992; Johar & Roggeveen, 2007). Repetition can even increase the perceived accuracy of plausible but false statements among participants who are subsequently able to identify the correct answer (Fazio et al., 2015).

Here we ask whether illusory truth effects extend to fake news. Given that the fake-news stories circulating on social media are quite different from the stimuli that have been employed in previous illusory truth experiments, in that they are implausible and highly partisan, finding such an effect for fake news extends the scope (and real-world relevance) of the illusory truth effect and, as we argue, informs theoretical models of the effect. Indeed, there are numerous reasons to think that simple prior exposure will not increase the perceived accuracy of fake news.

Implausibility as a Potential Boundary Condition of the Illusory Truth Effect

Fake-news stories are constructed with the goal of drawing attention and are therefore often quite fantastical and implausible. For example, Pennycook and Rand (2018a) gave participants a set of politically partisan fake-news headlines collected from online websites (e.g., ‘Trump to Ban All TV Shows That Promote Gay Activity Starting With Empire as President’) and found that they

were judged as accurate only 17.8% of the time. To contrast this figure with the existing illusory truth literature, Fazio et al. (2015) found that false trivia items were judged to be true around 40% of the time, even when restricting the analysis to participants who were subsequently able to recognize the statement as false. Thus, these previous statements (such as ‘chemosynthesis is the name of the process by which plants make their food’), despite being untrue, are much more plausible than are typical fake-news headlines. This may have consequences for whether repetition increases perceived accuracy of fake news: When it is completely obvious that a statement is false, it may be perceived as inaccurate regardless of how fluently it is processed. Although such an influence of plausibility is not explicitly part of the fluency-conditional model of illusory truth proposed by Fazio and colleagues (under which knowledge influences judgment only when people do not rely on fluency), the possibility of such an effect is acknowledged in their discussion when they state that they “expect that participants would draw on their knowledge, regardless of fluency, if statements contained implausible errors” (p. 1000). Similarly, when summarizing a meta-analysis of illusory truth effects, Dechêne et al. (2010) argued that “statements have to be ambiguous, that is, participants have to be uncertain about their truth status because otherwise the statements’ truthfulness will be judged on the basis of their knowledge” (p. 239). Thus, investigating the potential for an illusory truth effect for fake news is not simply important because it helps one to understand the spread of fake news but also because it allows one to test heretofore untested (but common) intuitions about the boundary conditions of the effect.

Motivated Reasoning as a Potential Boundary Condition of the Illusory Truth Effect

Another striking feature of fake news that may counteract the effect of repetition—and that is absent from prior studies of the illusory truth effect—is the fact that fake-news stories are not only political in nature but often extremely partisan. Although prior work has shown the illusory truth effect *on average* for (relatively innocuous) social-political opinion statements (Arkes, Hackett, & Boehm, 1989), the role of individual differences in ideological discordance has not been examined. Of importance, people have a strong motivation to reject the veracity of stories that conflict with their political ideology (Flynn, Nyhan, & Reifler, 2017; Kahan, 2013; Kahan et al., 2012), and the hyperpartisan nature of fake news makes such conflicts likely for roughly half the population. Furthermore, the fact that fake-news stories are typically of immediate real-world relevance—and therefore, presumably, more impactful on a person’s beliefs and actions than are the relatively trivial pieces of information considered in previous work on the illusory truth effect—should make people more inclined to think carefully about the accuracy of such stories, rather than rely on simple heuristics when making accuracy judgments. Thus, there is reason to expect that people may be resistant to illusory truth effects for partisan fake-news stories that they have politically motivated reasons to reject.

The Current Work

Although there are reasons why, in theory, people should not believe fake news (even if they have seen it before), it is clear that

many people do in fact find such stories credible. If repetition increases perceptions of accuracy for even highly implausible and partisan content, then increased exposure may (at least partly) explain why fake-news stories have recently proliferated. Here we assess this possibility with a set of highly powered and preregistered experiments. In a first study, we explored the impact of extreme implausibility on the illusory truth effect in the context of politically neutral statements. We found that implausibility does indeed present a boundary condition for illusory truth, such that repetition does not increase perceived accuracy of statements that essentially no one believes at baseline. In two more studies, however, we found that—despite being implausible, partisan, and provocative—fake-news headlines that are repeated are in fact perceived as more accurate. Taken together, these results shed light on how people come to have patently false beliefs, help to inform efforts to reduce such beliefs, and extend understanding of the basis of illusory truth effects.

Study 1: Extreme Implausibility Boundary Condition

Although existing models of the illusory truth effect do not explicitly take plausibility into account, we hypothesized that prior exposure should not increase perceptions of accuracy for statements that are *prima facie* implausible—that is, statements for which individuals hold extremely certain prior beliefs. In other words, when strong internal reasons exist to reject the veracity of a statement, it should not matter how fluently the statement is processed.

To assess implausibility as a boundary condition for the illusory truth effect, we created statements that participants would certainly know to be false (i.e., extremely implausible statements such as “The earth is a perfect square”) and manipulated prior exposure using a standard illusory truth paradigm (via Fazio et al., 2015). We also included unknown (but plausible) true and false trivia statements from a set of general knowledge norms (Tauber, Dunlosky, & Rawson, 2013). To balance out the set, we also gave participants obvious known truths (see Table 1 for example items from each set). Participants first rated the “interestingness” of half of the items, and following an unrelated intervening questionnaire, they were asked to assess the accuracy of all items. Thus, half of the items in the assessment stage were previously presented (i.e., familiarized), and half were novel. If implausibility is a boundary condition for the illusory truth effect, there should be no significant effect of repetition on extremely implausible (known) falsehoods. We expected to replicate the standard illusory truth effect for unknown (but plausible) trivia statements. For extremely plausible known true statements, there may be a ceiling effect on accuracy

judgments that precludes an effect of repetition (cf. results for fluency on known truths; Unkelbach, 2007).

Method

All data are available online (<https://osf.io/txf46/>). We preregistered our hypotheses, primary analyses, and sample size (<https://osf.io/txf46/>). Although one-tailed tests are justified in the case of preregistered directional hypotheses, here we followed conventional practices and used two-tailed tests throughout (the use of one-tailed vs. two-tailed tests does not qualitatively alter our results). All participants were recruited from Amazon’s Mechanical Turk (Horton, Rand, & Zeckhauser, 2011), which has been shown to be a reliable resource for research on political ideology (Coppock, 2016; Krupnikov & Levine, 2014; Mullinix, Leeper, Druckman, & Freese, 2015). These studies were approved by the Yale Human Subject Committee.

Participants. Our target sample was 500. In total, 566 participants completed some portion of the study. We had complete data for 515 participants (51 participants dropped out). Participants were removed if they indicated responding randomly ($N = 50$), searching online for any of the claims ($N = 24$; 1 of whom did not respond), or going through the familiarization stage without doing the task ($N = 32$). These exclusions were preregistered. The final sample ($N = 409$; mean age = 35.8 years) included 171 male and 235 female participants (three did not indicate their sex).

Materials. We created four known falsehoods (i.e., extremely implausible statements) and four known truths statements (see the online supplemental materials for a full list). We also used 10 true and 10 false trivia questions framed as statements (via Tauber et al., 2013). Trivia items were sampled from largely unknown facts (see Table 1).

Procedure. We used a procedure parallel to that used by Fazio et al. (2015). Participants were first asked to rate the “interestingness” of the items on a 6-point scale ranging from 1 (*very uninteresting*) to 6 (*very interesting*). Half of the items were presented in this familiarization stage (counterbalanced). Participants then completed a few demographic questions and the Positive and Negative Affect Schedule (PANAS; Watson, Clark, & Tellegen, 1988). This filler stage consisted of 25 questions and took approximately 2 min. Demographic questions consisted of age (“What is your age?”), sex (“What is your sex?”), education (“What is the highest level of school you have completed or the highest degree you have received,” with eight typical education-level options), English fluency (“Are you fluent in English?”), and zip code (“Please enter the ZIP code for your primary residence. Reminder: This survey is anonymous”). Finally, participants were asked to assess the accuracy of the statements on a 6-point scale ranging from 1 (*definitely false*) to 6 (*definitely true*). At the end of the survey, participants were asked about random responding (“Did you respond randomly at any point during the study?”) and use of search engines (“Did you search the Internet [via Google or otherwise] for any of the news headlines?”). Both were accompanied by a yes–no response option and the following clarification: “Note: Please be honest! You will get your HIT regardless of your response.”

Table 1
Example Items From Study 1, Which Vary as a Function of Whether they are Known/Unknown and True/False

| Type | Item |
|---------------------|--|
| Known | |
| True | There are more than fifty stars in the universe. |
| False (implausible) | The earth is a perfect square. |
| Unknown | |
| True | Billy the Kid’s last name was Bonney. |
| False | Angel Falls is located in Brazil. |

Results

Following our preregistration, the key comparison was between familiarized and novel implausible items. As predicted, repetition did not increase perceptions of accuracy for implausible (known false) statements ($p = .462$; see Table 2), whereas there was a significant effect of repetition for both true and false trivia (unknown) statements ($ps < .001$). There was no significant effect of repetition on very plausible (known true) statements ($p = .078$). These results were supported by a significant interaction between knowledge (known, unknown) and exposure (familiarized, novel), $F(1, 408) = 82.17$, $MSE = .35$, $p < .001$, $\eta^2 = .17$. Specifically, there was no significant overall effect of repetition for known items, $F(1, 408) = .91$, $MSE = .30$, $p = .341$, $\eta^2 = .002$, but a highly significant overall effect for unknown items, $F(1, 408) = 107.99$, $MSE = .47$, $p < .001$, $\eta^2 = .21$.

Discussion

Although we replicated prior results indicating a positive effect of repetition on ambiguously plausible statements, regardless of their correctness, we observed no significant effect of repetition on accuracy judgments for statements that were patently false.

Study 2: Fake News

Study 1 established that, at least, *extreme* implausibility is a boundary condition for the illusory truth effect. Nonetheless, given that fake-news stories are highly (but not entirely) implausible (Pennycook & Rand, 2017), it is unclear whether their level of plausibility would be sufficient to allow prior exposure to inflate the perceived accuracy of fake news. It is also unclear what impact the highly partisan nature of fake-news stimuli, and the motivated reasoning to which this partisanship may lead (i.e., reasoning biased toward conclusions that are concordant with previous opinion; Kahan, 2013; Kunda, 1990; Mercier & Sperber, 2011; Redlawsk, 2002), would have on any potential illusory truth effect. Motivated reasoning may cause people to see politically discordant stories as disproportionately inaccurate, such that the illusory truth effect may be diluted (or reversed) when headlines are discordant. We assessed these questions in Study 2.

In addition to assessing the baseline impact of repetition on fake news, we also investigated the impact of explicit warnings about a lack of veracity on the illusory truth effect, given that warnings have been shown to be effective tools for diminishing (although not abolishing) the memorial effects of misinformation (Ecker,

Lewandowsky, & Tang, 2010). Furthermore, such warnings are a key part of efforts to combat fake news—for example, Facebook's first major intervention against fake news consisted of flagging stories shown to be false with a caution symbol and the text *Disputed by 3rd Party Fact-Checkers* (Mosseri, 2016). To this end, half of the participants were randomly assigned to a warning condition in which this caution symbol and a *disputed* warning were applied to the fake-news headlines.

Prior work has shown that participants rate repeated trivia statements as more accurate than novel statements, even when they were told that the source was inaccurate (Begg et al., 1992). Specifically, Begg and colleagues (1992) attributed statements in the familiarization stage to people with either male or female names and then told participants that either all male or all female individuals were lying. Participants were then presented with repeated and novel statements—all without sources—and they rated previously presented statements as more accurate even if they had been attributed to the lying gender in the familiarization stage. This provides evidence that the illusory truth effect survives manipulations that decrease belief in statements at first exposure. Nonetheless, Begg and colleagues employed a design that was different in a variety of ways from our warning manipulation. Primarily, Begg and colleagues provided information about veracity indirectly: For any given statement presented during their familiarization phase, participants had to complete the additional step, at encoding, of mapping the source's gender into the information provided about which gender was unreliable in order to inform their initial judgment about accuracy. The *disputed* warnings we tested here, conversely, did not involve this extra mapping step. Thus, by assessing their impact on the illusory truth effect, we tested whether the scope of Begg and colleagues' findings extends to this more explicit warning, while also generating practically useful insight into the efficacy of this specific fake-news intervention.

Method

Participants. We had an original target sample of 500 participants in our preregistration. We then completed a full replication of the experiment with another 500 participants. Given the similarity across the two samples, the data sets were combined for the main analysis (the results are qualitatively similar when examining the two experiments separately; see the online supplemental materials). The first wave was completed on January 16, and the second wave was completed on February 3 (both in 2017). In total, 1,069 participants from Mechanical Turk completed some portion of the survey. However, 64 did not finish the study and were removed (33 from the no-warning condition and 31 from the warning condition). A further 32 participants indicated responding randomly at some point during the study and were removed. We also removed participants who reported searching for the headlines ($N = 18$) or skipping through the familiarization stage ($N = 6$). These exclusions were preregistered for Studies 1 and 3 but were accidentally omitted from the preregistration for Study 2. The results are qualitatively identical with the full sample, but we report analyses with participants removed to retain consistency across our studies. The final sample ($N = 949$; mean age = 37.1) included 449 male and 489 female participants (11 did not respond).

Table 2

Comparison of Familiarized and Novel Items for Known or Unknown True and False Statements in Study 1

| Type | Familiarized | Novel | Difference | <i>t</i> (df) | <i>p</i> |
|---------------------|--------------|-----------|------------|---------------|----------|
| Known | | | | | |
| True | 5.59 (.8) | 5.66 (.6) | -.07 | 1.77 (408) | .078 |
| False (implausible) | 1.13 (.6) | 1.11 (.5) | .02 | .74 (408) | .462 |
| Unknown | | | | | |
| True | 4.12 (.7) | 3.79 (.8) | .33 | 6.65 (408) | <.001 |
| False | 3.77 (.7) | 3.39 (.7) | .38 | 9.44 (408) | <.001 |

Note. Data presented are means, with standard deviations in parentheses.




Figure 1. Sample fake-news headline without (Panel A) and with (Panel B) a *Disputed by 3rd Party Fact-Checkers* warning, as presented in Studies 2 and 3. The original image (but not the headline) has been replaced with a stock military image (under a CC0 license) for copyright purposes. Image is from <https://www.pexels.com/photo/flight-sky-sunset-men-54098/>. See the online article for the color version of this figure.

Materials and procedure. Participants engaged in a three-stage experiment. In the familiarization stage, participants were shown six news headlines that were factually accurate (real news) and six others that were entirely untrue (fake news). The headlines were presented in a format identical to that of Facebook posts (i.e., a headline with an associated photograph above it and a lede sentence and byline below it; see Figure 1A; fake and real-news headlines can be found in the Appendix—images for each item, as presented to participants, can be found at the following link: <https://osf.io/txf46/>). Participants were randomized into two conditions: (a) The warning condition, where all of the fake-news headlines (but none of the real-news headlines) in the familiarization stage were accompanied by a *Disputed by 3rd Party Fact-Checkers* tag” (see Figure 1B), or (b) the control condition, where fake and real-news headlines were displayed without warnings. In the familiarization stage, participants engaged with the news headlines in an ecologically valid way: They indicated whether they would share each headline on social media. Specifically, participants were asked “Would you consider sharing this story online (for example, through Facebook or Twitter)?” and were given three response options (*No, Maybe, Yes*). For purposes of data analysis, *No* was coded as 0 and *Maybe* and *Yes* were coded as 1.¹

The participants then advanced to the distractor stage, in which they completed a set of filler demographic questions. These included age, sex, education, proficiency in English, political party (Democratic, Republican, Independent, other), social and economic conservatism (separate items),² and two questions about the 2016 election. For these election-related questions, participants were first asked to indicate who they voted for (given the following options: *Hillary Clinton, Donald Trump, Other Candidate [such as Jill Stein or Gary Johnson], I did not vote for reasons outside my control, I did not vote but I could have, and I did not vote out of protest*). Participants were then asked “If you absolutely

had to choose between only Clinton and Trump, who would you prefer to be the next President of the United States.” This binary response was then used as our political ideology variable for the concordance–discordance analysis. Specifically, for participants who indicated a preference for Trump, pro-Republican stories were scored as politically concordant and pro-Democrat stories were scored as politically discordant; for participants who indicated a preference for Clinton, pro-Democrat stories were scored as politically concordant and pro-Republican stories were scored as politically discordant. The filler stage took approximately 1 min.

Finally, participants entered the assessment stage, where they were presented with 24 news headlines—the 12 headlines they saw in the familiarization stage and 12 new headlines (six fake news, six real news)—and rated each for familiarity and accuracy. Which headlines were presented in the familiarization stage was counter-balanced across participants, and headline order was randomized for every participant in both Stage 1 and Stage 3. Moreover, the items were balanced politically, with half being pro-Democrat and half pro-Republican. The fake-news headlines were selected from Snopes.com, a third-party website that fact-checks news stories. The real headlines were contemporary stories from mainstream news outlets. For each item, participants were first asked “Have you seen or heard about this story before?” and were given three

¹ This was not preregistered for Study 2; however, it was for Study 3. Hence, we used this analysis strategy to retain consistency across the two fake-news studies. The results are qualitatively similar if the social media question is scored continuously.

² Participants answered the prompt “On social issues I am” with *Strongly Liberal, Somewhat Liberal, Moderate, Somewhat Conservative, or Strongly Conservative*. The same was true for the economic conservatism item except the prompt was “On economic issues I am.”

response options (*No*, *Unsure*, *Yes*). For the purposes of data analysis, *No* and *Unsure* were combined (this was preregistered in Study 3 but not in Study 2). As in other work on perceptions of news accuracy (Pennycook & Rand, 2017, 2018a, 2018b), participants were then asked “To the best of your knowledge, how accurate is the claim in the above headline?” and they rated accuracy on the following 4-point scale: 1 (*not at all accurate*), 2 (*not very accurate*), 3 (*somewhat accurate*), and 4 (*very accurate*). We focused on judgments about news headlines, as opposed to full articles, because much of the public’s engagement with news on social media involves reading only story headlines (Gabiellov, Ramachandran, Chaintreau, & Legout, 2016).

At the end of the survey, participants were asked about random responding, use of search engines to check accuracy of the stimuli, and whether they skipped through the familiarization stage (“At the beginning of the survey [when you were asked whether you would share the stories on social media], did you just skip through without reading the headlines?”). All were accompanied by a yes–no response option.

Our preregistration specified the comparison between familiarized and novel fake news, separately in the warning and no-warning conditions, as the key analyses. However, for completeness, we report the full set of analyses that emerge from our mixed-design analysis of variance (ANOVA). Our political concordance analysis deviates somewhat from the analysis that was preregistered, and our follow-up analysis that focuses on unfamiliar headlines was not preregistered. Our full preregistration is available at the following link: <https://osf.io/txf46/>.

Results

As a manipulation check for our familiarization procedure, we submitted familiarity ratings (recorded during the assessment stage) to a 2 (type: fake, real) \times 2 (exposure: familiarized, novel) \times 2 (warning: warning, no-warning) mixed-design ANOVA. Critically, there was a main effect of exposure such that familiarized headlines were rated as more familiar ($M = 44.7\%$, $SD = 35.6$) than were novel headlines ($M = 16.2\%$, $SD = 15.5$), $F(1, 947) = 578.76$, $MSE = .13$, $p < .001$, $\eta^2 = .38$, and a significant simple effect was present within every combination of news type and warning condition (all t s > 14.0 , all p s $< .001$). This indicates that our social media sharing task in the familiarization stage was sufficient to capture participants’ attention (further analysis of familiarity judgments can be found in the online supplemental materials).

As a manipulation check for attentiveness to the *Disputed by 3rd Party Fact-Checkers* warning, we submitted the willingness to share news articles on social media measure (from the familiarization stage) to a 2 (type: fake, real) \times 2 (condition: warning, no-warning) mixed-design ANOVA. This analysis revealed a significant main effect of type, such that our participants were more willing to share real stories ($M = 41.6\%$, $SD = 31.8$) than fake stories ($M = 29.7\%$, $SD = 29.8$), $F(1, 947) = 131.16$, $MSE = .05$, $p < .0017$, $\eta^2 = .12$. More important, there was a significant main effect of condition, $F(1, 947) = 15.33$, $MSE = .13$, $p < .001$, $\eta^2 = .016$, which was qualified by an interaction between type and condition, $F(1, 947) = 19.65$, $MSE = .05$, $p < .001$, $\eta^2 = .020$, such that relative to the no-warning condition, participants in the warning condition reported being less willing to share fake-news

headlines (which actually bore the warnings in the warning condition; warning: $M = 23.9\%$, $SD = 28.3$; no-warning: $M = 35.2\%$, $SD = 30.2$), $t(947) = 5.93$, $p < .001$, $d = .39$, whereas there was no significant difference across conditions in sharing of real news (which did not have warnings in either condition; warning: $M = 40.6\%$, $SD = 32.2$; no-warning: $M = 42.6\%$, $SD = 31.5$; $t < 1$). Thus, participants clearly paid attention to the warnings.

We now turn to perceived accuracy, our main focus. Perceived accuracy was entered into a 2 (type: fake, real) \times 2 (exposure: familiarized, novel) \times 2 (warning: warning, no-warning) mixed-design ANOVA (see Table 3 for means and standard deviations). Demonstrating the presence of an illusory truth effect, there was a significant main effect of exposure, $F(1, 947) = 93.65$, $MSE = .12$, $p < .001$, $\eta^2 = .09$, such that headlines presented in the familiarization stage ($M = 2.24$, $SD = .42$) were rated as more accurate than were novel headlines ($M = 2.13$, $SD = .39$). There was also a significant main effect of headline type, such that real-news headlines ($M = 2.67$, $SD = .48$) were rated as much more accurate than were fake-news headlines ($M = 1.71$, $SD = .46$), $F(1, 945) = 2,424.56$, $MSE = .36$, $p < .001$, $\eta^2 = .72$. However, there was no significant interaction between exposure and type of news headline ($F < 1$). In particular, prior exposure increased accuracy ratings even when considering only fake-news headlines (see Figure 2; familiarized: $M = 1.77$, $SD = .56$; novel: $M = 1.65$, $SD = .48$), $t(948) = 7.60$, $p < .001$, $d = .25$. For example, nearly twice as many participants (92.1% increase, from 38 to 73 out of 949 total) judged the fake-news headlines presented to them during the familiarization stage as accurate (mean accuracy rating above 2.5), compared to the stories presented to them for the first time in the assessment stage. Although both of these participant counts are only a small fraction of the total sample, the fact that a single exposure to the fake stories doubled the number of credulous participants suggests that repetition effects may have a substantial impact in daily life, where people can see fake-news headlines cycling many times through their social media newsfeeds.

What effect did the presence of warnings on fake news in the familiarization stage have on later judgments of accuracy and, potentially, the effect of repetition? The ANOVA just described revealed a significant main effect of the warning manipulation, $F(1, 947) = 5.39$, $MSE = .53$, $p = .020$, $\eta^2 = .005$, indicating that the warning decreased perceptions of news accuracy. However, this was qualified by an interaction between warning and type, $F(1, 947) = 5.83$, $MSE = .36$, $p = .016$, $\eta^2 = .006$. Whereas the presence of warnings on fake news in the assessment stage had no effect on perceptions of real-news accuracy (warning: $M = 2.67$, $SD = .49$; no-warning: $M = 2.67$, $SD = .48$; $t < 1$), participants rated fake news as less accurate in the warning condition (warning: $M = 1.66$, $SD = .46$; no-warning: $M = 1.76$, $SD = .46$), $t(947) = 3.40$, $p = .001$, $d = .22$. Furthermore, there was a marginally significant interaction between exposure and warning, $F(1, 947) = 3.32$, $MSE = .12$, $p = .069$, $\eta^2 = .004$, such that the decrease in overall perceptions of accuracy was significant for familiarized items (warning: $M = 2.21$, $SD = .41$; no-warning: $M = 2.28$, $SD = .43$), $t(947) = 2.77$, $p = .006$, $d = .18$, but not novel items, (warning: $M = 2.12$, $SD = .38$; no-warning: $M = 2.15$, $SD = .39$), $t(947) = 1.36$, $p = .175$, $d = .09$. That is, the warning decreased perceptions of accuracy for items that were presented in the familiarization stage—both fake stories that were labeled with

Table 3
Comparison of Familiarized and Novel Items for Politically Concordant and Discordant Items in the Warning and No-Warning Conditions

| Type and warning status | Familiarized | Novel | <i>t</i> (df) | <i>p</i> |
|-------------------------|--------------|-----------|---------------|----------|
| Politically concordant | | | | |
| Fake news | | | | |
| No Warning | 1.93 (.7) | 1.78 (.6) | 5.46 (486) | <.001 |
| Warning | 1.81 (.7) | 1.68 (.6) | 4.69 (459) | <.001 |
| Real news | | | | |
| No Warning | 2.98 (.6) | 2.83 (.7) | 5.45 (486) | <.001 |
| Warning | 2.92 (.7) | 2.86 (.7) | 2.17 (459) | .031 |
| Politically discordant | | | | |
| Fake news | | | | |
| No Warning | 1.72 (.6) | 1.60 (.5) | 3.91 (486) | <.001 |
| Warning | 1.60 (.6) | 1.53 (.5) | 2.66 (459) | .008 |
| Real news | | | | |
| No Warning | 2.50 (.6) | 2.39 (.6) | 3.85 (486) | <.001 |
| Warning | 2.49 (.6) | 2.40 (.6) | 3.03 (459) | .003 |

Note. Data presented are means, with standard deviations in parentheses. Politically concordant items consisted of pro-Democrat items for Clinton supporters and pro-Republican items for Trump supporters (and vice versa for politically discordant items).

warnings and the real stories presented without warnings (see footnote 3)—but not for items that were not presented in the familiarization stage.

There was no significant three-way interaction, however, between headline type, exposure, and warning condition ($F < 1$). As a consequence, the repetition effect was evident for fake-news headlines in the warning condition, $t(460) = 4.89$, $p < .001$, $d = .23$, as well as in the no-warning condition, $t(487) = 5.81$, $p < .001$, $d = .26$ (see Figure 2). That is, participants rated familiarized fake-news headlines that they were explicitly warned about as more accurate were than novel fake-news headlines that they were not warned about (despite the significant negative effect of warnings on perceived accuracy of fake news reported earlier). In fact, there was no significant interaction between the exposure and warning manipulations when isolating the analysis to fake-news headlines, $F(1, 947) = 1.00$, $MSE = .12$, $p = .317$, $\eta^2 = .001$. Thus, the warning seems to have created a general sense of distrust—thereby reducing perceived accuracy for both familiarized and novel fake-news headlines—rather than making people particularly distrust the stories that were labeled as disputed.

As a secondary analysis,³ we also investigated whether the effect of prior exposure is robust to political concordance (i.e., whether headlines were congruent or incongruent with one's political stance). Mean perceptions of news accuracy for politically concordant and discordant items as a function of type, exposure, and warning condition can be found in Table 3. Perceived accuracy was entered into a 2 (political valence: concordant, discordant) \times 2 (type: fake, real) \times 2 (exposure: familiarized, novel) \times 2 (warning: warning, no-warning) mixed-design ANOVA. First, as a manipulation check, politically concordant items were rated as far more accurate than were politically discordant items overall, $F(1, 945) = 573.08$, $MSE = .34$, $p < .001$, $\eta^2 = .38$ (see Table 3). Nonetheless, we observed no significant interaction between the repetition manipulation and political valence, $F(1, 945) = 2.24$,

$MSE = .15$, $p = .135$, $\eta^2 = .002$. The illusory truth effect was evident for fake-news headlines that were politically discordant, $t(946) = 4.70$, $p < .001$, $d = .15$, as well as concordant, $t(946) = 7.19$, $p < .001$, $d = .23$. Political concordance interacted significantly with type of news story, $F(1, 945) = 138.91$, $MSE = .23$, $p < .001$, $\eta^2 = .13$, such that the difference between perceptions of real and fake news (i.e., discernment) was greater for politically concordant headlines (real news: $M = 2.90$, $SD = .59$; fake news: $M = 1.80$, $SD = .56$), than for politically discordant headlines (real news: $M = 2.44$, $SD = .53$; fake news: $M = 1.61$, $SD = .48$), $t(946) = 11.8$, $p < .001$, $d = .38$ (see Pennycook & Rand, 2018a, for a similar result). All other interactions with political concordance were not significant (all F s < 1.5 , $ps > .225$).

The illusory truth effect also persisted when analyzing only news headlines that the participants marked as unfamiliar (i.e., in the same mixed-design ANOVA as mentioned earlier but analyzing only stories the participants were not consciously aware of having seen in the familiarization stage or at some point prior to the experiment; familiarized: $M = 1.90$, $SD = .53$; novel: $M = 1.83$, $SD = .49$), $F(1, 541)^4 = 11.82$, $MSE = .17$, $p = .001$, $\eta^2 = .02$ (see the online supplemental materials for details and further statistical analysis).

Discussion

The results of Study 2 indicate that a single prior exposure is sufficient to increase perceived accuracy for both fake and real news. This occurs even (a) when fake news is labeled as *Disputed by 3rd Party Fact-Checkers* during the familiarization stage (i.e., during encoding at first exposure), (b) among fake (and real) news headlines that are inconsistent with one's political ideology, and (c) when isolating the analysis to news headlines that participants were not consciously aware of having seen in the familiarization stage.

Study 3: Fake News, 1-Week Interval

We next sought to assess the robustness of our finding that repetition increases perceptions of fake-news accuracy by making two important changes to the design of Study 2. First, we assessed the persistence of the repetition effect by inviting participants back after a weeklong delay (following previous research that has shown illusory truth effects to persist over substantial periods of time; e.g., Hasher et al., 1977; Schwartz, 1982). Second, we restricted our analyses to only those items that were unfamiliar to participants when entering the study, which allows for a cleaner novel baseline.

³ These analyses were not preregistered, although we did preregister a parallel analysis where pro-Democrat and pro-Republican items would be analyzed separately while comparing liberals and conservatives. The present analysis simply combines the data into a more straightforward analysis and uses the binary Clinton–Trump choice to distinguish liberals and conservatives. The effect of prior exposure was significant for fake news when political concordance was determined based on Democrat–Republican party affiliation: politically concordant, $t(609) = 4.8$, $p < .001$; politically discordant, $t(609) = 2.9$, $p = .004$.

⁴ Degrees of freedom are lower here because this analysis includes only individuals who were unfamiliar with at least one item in each cell of the design (familiarized–novel and fake news–real news).




Figure 2. Exposing participants to fake-news headlines in Study 2 increased accuracy ratings, even when the stories were tagged with a warning indicating that they had been disputed by third-party fact-checkers. Panel A: Mean accuracy ratings for fake-news headlines as a function of repetition (familiarized stories were shown previously during the familiarization stage; novel stories were shown for the first time during the assessment stage) and presence or absence of a warning that fake-news headlines had been disputed. Error bars indicate 95% confidence intervals. Panel B: Distribution of participant-average accuracy ratings for the fake-news headlines, comparing the six familiarized stories shown during the familiarization stage (red; lower “mountain” for print version readers) with the six novel stories shown for the first time in the assessment stage (blue; the higher “mountain”). We collapsed across warning and no-warning conditions because the repetition effect did not differ significantly by condition. See the online article for the color version of this figure.

Method

Participants. Our target sample was 1,000 participants from Mechanical Turk. This study was completed on February 1 and 2, 2017. Participants who completed Study 2 were not permitted to complete Study 3. In total, 1,032 participants completed the study, 40 of which dropped out or had missing data (14 from the no-warning condition, 26 from the warning condition). Participants who reported responding randomly ($N = 29$), skipping over the familiarization phase ($N = 1$), or searching online for the headlines ($N = 22$) were removed. These exclusions were preregistered. The final sample ($N = 940$; mean age = 36.8) included 436 male and 499 female participants (five did not respond).

Materials and procedure. The design was identical to that of Study 2 (including the warning and no-warning conditions), with a few exceptions. First, the length of the distractor stage was increased by adding 20 unrelated questionnaire items to the demographics questions (namely, the PANAS, as in Study 1). This filler stage took approximately 2 min to complete. Furthermore, participants were invited to return for a follow-up session 1 week later in which they were presented with the same headlines they had seen in the assessment stage plus a set of novel headlines not included in Session 1 ($N = 566$ participants responded to the follow-up invitation). To allow full counterbalancing, we presented participants with eight headlines in the familiarization phase, 16 headlines in the accuracy judgment phase (of which eight were those shown in the familiarization phase), and 24 headlines in the follow-up session a week later (of which 16 were those shown in the assessment phase of Session 1), again maintaining an equal number of real-fake and pro-Democrat-pro-Republican headlines within each block. The design of Study 3 therefore allowed us to assess the temporal stability of the repetition effect within both Session 1 (over the span of a distractor task) and Session 2 (over the span of a week).

Second, during the familiarization stage participants were asked to indicate whether each headline was familiar, instead of whether

they would share the story on social media (the social media question was moved to the assessment stage). This modification allowed us to restrict our analyses to only those items that were unfamiliar to participants when entering the study (i.e., they said “no” when asked about familiarity),⁵ allowing for a cleaner assessment of the causal effect of repetition (903 of the 940 participants in Session 1 were previously unfamiliar with at least one story of each type and were thus included in the main text analysis, as were 527 out of the 566 participants in Session 2; see the online supplemental materials for analyses with all items and all participants). Fake- and real-news headlines as presented to participants can be found at the following link: <https://osf.io/txf46/>.

As in Experiment 2, our preregistration specified the comparison between familiarized and novel fake news in both the warning and no-warning conditions (and for both sessions) as the key analyses, although in this case we preregistered the full 2 (type: fake, real) $\times 2$ (exposure: familiarized, novel) $\times 2$ (warning: warning, no-warning) mixed-design ANOVA. We also preregistered the political concordance analysis. Finally, we preregistered the removal of cases where participants were familiar with the news headlines as a secondary analysis, but we focus on it as a primary analysis here because this is the novel feature relative to the case in Study 2 (primary analyses including all participants are discussed in footnote 8). Our preregistration is available at the following link: <https://osf.io/txf46/>.

⁵ Whereas participants indicated their familiarity with familiarized items prior to completing the accuracy judgments, they indicated their familiarity with novel items after completing the accuracy judgments. Thus, it is possible that seeing the news headlines in the accuracy-judgment phase would increase perceived familiarity. There was no evidence for this, however, because mean familiarity judgment (scored continuously) did not significantly differ based on whether the judgment was made before or after the accuracy judgment phase, $t(939) = .68$, $SE = .01$, $p = .494$, $d = .02$. Participants were unfamiliar with 81.2% of the fake-news headlines and 49.2% of the real-news headlines.

Results

Perceived accuracy was entered into a 2 (type: fake, real) \times 2 (exposure: familiarized, novel) \times 2 (warning: warning, no-warning) mixed-design ANOVA. Replicating the illusory truth effect from Study 2, there was a clear causal effect of prior exposure on accuracy in Session 1 of Study 3 despite the longer distractor stage: Headlines presented in the familiarization stage ($M = 2.01$, $SD = .54$) were rated as more accurate than were novel headlines ($M = 1.92$, $SD = .49$), $F(1, 721) = 22.52$, $MSE = .23$, $p < .001$, $\eta^2 = .03$. Again replicating the results of Study 2, there was a significant main effect of type, such that real stories ($M = 2.31$, $SD = .63$) were rated as much more accurate than were fake stories ($M = 1.63$, $SD = .52$), $F(1, 721) = 934.57$, $MSE = .36$, $p < .001$, $\eta^2 = .57$, but there was no significant interaction between exposure and type of news headline, $F(1, 721) = 2.65$, $MSE = .20$, $p = .104$, $\eta^2 = .004$. Accordingly, prior exposure increased perceived accuracy even when considering only fake-news headlines (see Figure 3A and 3B), $t(902)^6 = 5.99$, $p < .001$, $d = .20$ (89.5% increase in number of participants judging familiarized fake-news headlines as accurate compared to novel fake-news headlines; i.e., from 38 to 72 participants out of 903).

Unlike the case in Study 2, there was no main effect of the warning manipulation on overall perceptions of accuracy (i.e., across the aggregate of fake and real news; $F < 1$). However, there was a marginally significant interaction between type of news story and warning condition, $F(1, 721) = 2.95$, $MSE = .36$, $p = .086$, $\eta^2 = .004$. Regardless, the fake-news warnings in the familiarization stage had no significant overall effect on perceptions of fake-news accuracy in the assessment stage (warning: $M = 1.61$, $SD = .50$; no-warning: $M = 1.66$, $SD = .54$), $t(932)^7 = 1.54$, $p = .123$, $d = .10$. There was also no significant effect of the warning on perceptions of real-news accuracy (warning: $M = 2.32$, $SD = .63$; no-warning: $M = 2.30$, $SD = .63$; $t < 1$), no significant interaction between the repetition and warning manipulations, $F(1, 721) = 1.89$, $MSE = .23$, $p = .169$, $\eta^2 = .003$, and no significant three-way interaction between warning, exposure, and type of news story ($F < 1$).⁸ Nonetheless, it should be noted that familiarized fake-news headlines (i.e., the fake-news headlines that were warned about in the familiarization stage) were rated as less accurate ($M = 1.64$, $SD = .59$) than were the same headlines in the control (no-warning) condition ($M = 1.73$, $SD = .63$), $t(925) = 2.14$, $p = .032$, $d = .14$, suggesting that the warning did have some effect on accuracy judgments. However, this effect was smaller than in Study 2 and did not extend to nonwarned (and not familiarized) fake news. This is perhaps due to the smaller number of items in the familiarization stage of Study 3.

Following our preregistration, we also analyzed the effect of exposure for fake-news headlines separately in the warning and no-warning conditions. The repetition effect was evident for fake-news headlines in both the warning condition (familiarized: $M = 1.63$, $SD = .58$; novel: $M = 1.55$, $SD = .52$), $t(447) = 3.07$, $p = .002$, $d = .14$, and the no-warning condition (familiarized: $M = 1.71$, $SD = .61$; novel: $M = 1.58$, $SD = .54$), $t(454) = 5.41$, $p < .001$, $d = .25$. Furthermore, familiarized fake-news headlines were judged as more accurate than were novel ones for both political discordant items (familiarized: $M = 1.60$, $SD = .67$; novel: $M = 1.51$, $SD = .63$), $t(858) = 3.41$, $p = .001$, $d = .12$, and political concordant items

(familiarized: $M = 1.72$, $SD = .77$; novel: $M = 1.59$, $SD = .67$), $t(801) = 4.93$, $p < .001$, $d = .18$; note that an ANOVA including concordance indicated that there was no significant interaction between repetition and political concordance for fake news, $F(1, 769) = 1.46$, $MSE = .32$, $p = .228$, $\eta^2 = .002$.⁹

Following up 1 week later, we continued to find a clear causal effect of repetition on accuracy ratings: Perceived accuracy of a story increased linearly with the number of times the participants had been exposed to that story. Using linear regression with robust standard errors clustered on participant,¹⁰ we found a significant positive relationship between number of exposures and accuracy overall (familiarized twice: $M = 2.00$, $SD = .53$; familiarized once: $M = 1.94$, $SD = .53$; novel: $M = 1.90$, $SD = .51$; $b = .046$), $t(537) = 3.68$, $p < .001$, and when considering only fake-news headlines (see Figure 3C; familiarized twice: $M = 1.70$, $SD = .58$; familiarized once: $M = 1.66$, $SD = .58$; novel: $M = 1.60$, $SD = .53$; $b = .048$), $t(526) = 3.66$, $p < .001$ (64% increase in number of participants judging fake-news headlines as accurate among stories seen twice compared to novel fake-news headlines; i.e., from 25 to 41 participants out of 527). Once again, this relationship was evident for fake news in both the warning condition (familiarized twice: $M = 1.67$, $SD = .59$; familiarized once: $M = 1.63$, $SD = .56$; novel: $M = 1.60$, $SD = .52$; $b = .036$), $t(276) = 1.97$, $p = .050$, and the no-warning condition (familiarized twice: $M = 1.73$, $SD = .57$; familiarized once: $M = 1.70$, $SD = .59$; novel: $M = 1.61$, $SD = .53$; $b = .061$), $t(249) = 3.27$, $p = .001$;

⁶ Only unfamiliar headlines are included, and therefore missing data account for missing participants in some cells of the design. Degrees of freedom vary throughout because the maximum number of participants is included in each analysis.

⁷ Degrees of freedom change here because this analysis includes the maximum number of individuals who were unfamiliar with at least one fake-news item.

⁸ In our (also) preregistered analysis that includes both previously familiar and unfamiliar items, there is a main effect of repetition, $F(1, 938) = 18.98$, $MSE = .16$, $p < .001$, $\eta^2 = .02$, but (unlike in Study 2) a significant interaction between exposure and warning condition, $F(1, 938) = 7.81$, $MSE = .16$, $p = .005$, $\eta^2 = .01$. There was a significant repetition effect for fake news in the no-warning condition, $t(475) = 5.31$, $SE = .03$, $p < .001$, $d = .24$, but no effect in the warning condition, $t(463) = 1.30$, $SE = .03$, $p = .193$, $d = .06$. It is possible that prior knowledge of the items facilitated explicit recall of the warning, which may have mitigated the illusory truth effect (see the online supplemental materials for means and further analyses).

⁹ We focused on fake-news headlines here because the political concordance manipulation cuts the number of items in half. Including real news in this analysis decreases the number of participants markedly because the analysis of variance (ANOVA) requires each participant to contribute at least one observation to each cell of the design. Nonetheless, the full ANOVA reveals a significant main effect of repetition, $F(1, 312) = 8.94$, $p = .003$, $\eta^2 = .03$, and no interaction with political concordance ($F < 1$). The effect of prior exposure was also significant for fake news when political concordance was determined based on Democrat–Republican party affiliation: politically concordant, $t(494) = 4.1$, $p < .001$; politically discordant, $t(529) = 2.3$, $p = .020$.

¹⁰ This specific analysis was not preregistered. Rather, the preregistration called for a comparison of the full 16 items from Session 1 with the eight novel items in Session 2. This, too, revealed a significant main effect of repetition (using the same analysis of variance as in the Session 1 analysis), $F(1, 453) = 12.91$, $p < .001$, $\eta^2 = .03$. However, such an analysis does not illuminate the increasing effect of exposure, hence our deviation from the preregistration (see the online supplemental materials for further details and analyses).

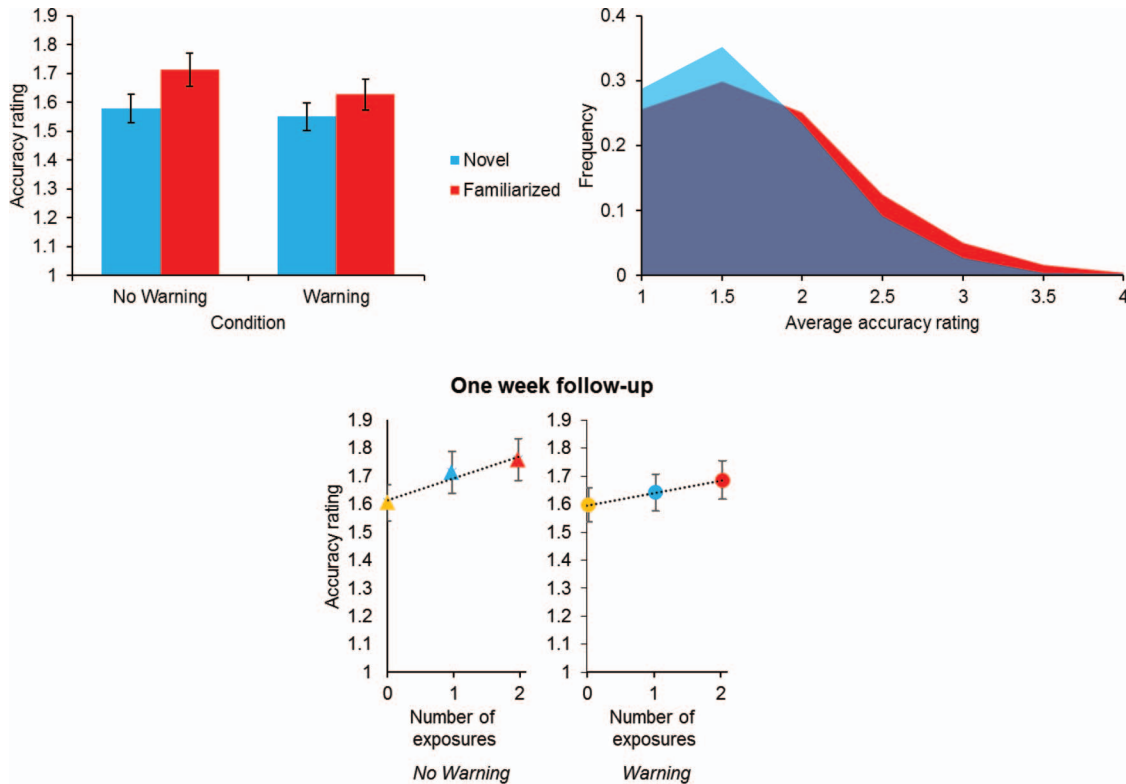


Figure 3. The illusory truth effect for fake news is persistent, lasting over a longer filler stage in Study 3 and continuing to be observed in a follow-up session 1 week later. Panel A: Mean accuracy ratings for fake-news headlines in Session 1 of Study 3 as a function of repetition and presence or absence of a warning that fake-news headlines had been disputed. Error bars indicate 95% confidence intervals. Panel B: Distribution of participant-average accuracy ratings for the fake-news headlines in Study 3, comparing the four headlines shown during the familiarization stage (red; the lower “mountain” for print version readers) with the four novel headlines shown for the first time in the assessment stage (blue; the higher “mountain”). We collapsed across warning and no-warning conditions because the repetition effect did not differ significantly by condition. Panel C: Mean accuracy ratings for fake-news headlines in the follow-up session conducted 1 week later, as a function of number of exposures to the story (two times for headlines previously presented in the familiarization and assessment stage of Session 1; one time for headlines previously presented in only the assessment stage of Session 1; and no times for headlines introduced for the first time in the follow-up session) and presence or absence of warning tag. Error bars indicate 95% confidence intervals based on robust standard errors clustered by participant, and the trend line is shown in dotted black. See the online article for the color version of this figure.

note that there was no significant interaction between the repetition and warning manipulations ($b = -.025$), $t(526) = .96$, $p = .337$ (see Figure 3C). This relationship was also evident for fake-news headlines that were politically discordant (familiarized twice: $M = 1.62$, $SD = .72$; familiarized once: $M = 1.61$, $SD = .68$; novel: $M = 1.54$, $SD = .62$; $b = .041$), $t(525) = 2.28$, $p = .023$, as well as politically concordant (familiarized twice: $M = 1.78$, $SD = .77$; familiarized once: $M = 1.71$, $SD = .75$; novel: $M = 1.66$, $SD = .70$; $b = .061$), $t(523) = 3.24$, $p = .001$.

Discussion

The results of Study 3 further demonstrated that prior exposure increases perceived accuracy of fake news. This occurred regardless of political discordance and among previously unfamiliar headlines that were explicitly warned about during familiarization.

Crucially, the effect of repetition on perceived accuracy persisted after a week and increased with an additional repetition. This suggests that fake-news credulity compounds with increasing exposures and maintains over time.

General Discussion

Although repetition did not impact accuracy judgments of totally implausible statements, across two preregistered experiments with a total of more than 1,800 participants we found consistent evidence that repetition did increase the perceived accuracy of fake-news headlines. Indeed, a single prior exposure to fake-news headlines was sufficient to measurably increase subsequent perceptions of their accuracy. Although this effect was relatively small ($d = .20-.21$), it increased with a second exposure, thereby suggesting a compounding effect of repetition across time. Explicit-

itly warning individuals that the fake-news headlines had been disputed by third-party fact-checkers (which was true in every case) did not abolish or even significantly diminish this effect. Furthermore, the illusory truth effect was evident even among news headlines that were inconsistent with the participants' stated political ideology.

Mechanisms of Illusory Truth

First, it is important to note that repetition increased accuracy even for items that the participants were not consciously aware of having been exposed to. This supports the broad consensus that repetition influences accuracy through a low-level fluency heuristic (Alter & Oppenheimer, 2009; Begg et al., 1992; Reber et al., 1998; Unkelbach, 2007; Whittlesea, 1993). These findings indicate that our repetition effect was likely driven, at least in part, by automatic (as opposed to strategic) memory retrieval (Diana, Yonelinas, & Ranganath, 2007; Yonelinas, 2002; Yonelinas & Jacoby, 2012). More broadly, these effects correspond with prior work demonstrating the power of fluency to influence a variety of judgments (Schwarz, Sanna, Skurnik, & Yoon, 2007); for example, subliminal exposure to a variety of stimuli (e.g., Chinese characters) increases associated positive feelings (i.e., the mere exposure effect; see Zajonc, 1968, 2001). Our evidence that the illusory truth effect extends to implausible and even politically inconsistent fake-news stories expands the scope of these effects. That perceptions of fake-news accuracy can be manipulated so easily despite being highly implausible (only 15%–22% of the headlines were judged to be accurate) has substantial practical implications (discussed later). However, what implications do these results have for the understanding of the mechanisms that underlie the illusory truth effect (and, potentially, a broader array of fluency effects observed in the literature)?

For decades, it has been assumed that repetition increases accuracy for only statements that are ambiguous (e.g., Dechêne et al., 2010) because, otherwise, individuals will simply use prior knowledge to determine truth. However, recent evidence has indicated that repetition can increase the perceived accuracy of even plausible but false statements (e.g., 'chemosynthesis is the name of the process by which plants make their food') among participants who were subsequently able to identify the correct answer (Fazio et al., 2015). However, it may be that the illusory truth effect is robust to the presence of conflicting prior knowledge only when statements are plausible enough that individuals fail to detect the conflict (for a perspective on conflict detection during reasoning, see Pennycook, Fugelsang, & Koehler, 2015). Indeed, as noted earlier, Fazio and colleagues (2015) speculated that "participants would draw on their knowledge, regardless of fluency, if statements contained implausible errors" (p. 1000). On the contrary, our findings indicate that implausibility is only a boundary condition of the illusory truth effect in the extreme: It is possible to use repetition to increase the perceived accuracy even for entirely fabricated and, frankly, outlandish fake-news stories that, given some reflection (Pennycook & Rand, 2018a, 2018b), people probably know are untrue. This observation substantially expands the purview of the illusory truth effect and suggests that external reasons for disbelief (such as direct prior knowledge

and implausibility) are no safeguard against the fluency heuristic.

Motivated Reasoning

Our results also have implications for a broad debate about the scope of motivated reasoning, which has been taken to be a fundamental aspect of how individuals interact with political misinformation and disinformation (Swire, Berinsky, Lewandowsky, & Ecker, 2017) and has been used to explain the spread of fake news (Allcott & Gentzkow, 2017; Beck, 2017; Calvert, 2017; Kahan, 2017; Singal, 2017). Although Trump supporters were indeed more skeptical about fake-news headlines that were anti-Trump relative to Clinton supporters (and vice versa), our results show that repetition increases perceptions of accuracy even in such politically discordant cases. Take, for example, the item "BLM Thug Protests President Trump With Selfie . . . Accidentally Shoots Himself in the Face," which is politically discordant for Clinton supporters and politically concordant for Trump supporters. Whereas on first exposure Clinton supporters were less likely (11.7%) than Trump supporters (18.5%) to rate this headline as accurate, suggesting the potential for motivated reasoning, a single prior exposure to this headline increased accuracy judgments in both cases (to 17.9% and 35.5%, for Clinton and Trump supporters, respectively). Thus, fake-news headlines were positively affected by repetition even when there was a strong political motivation to reject them. This observation complements the results of Pennycook and Rand (2018a), who found—in contrast to common motivated reasoning accounts (Kahan, 2017)—that analytic thinking leads to disbelief in fake news regardless of political concordance. Taken together, this suggests that motivated reasoning may play less of a role in the belief in fake news than is often argued.

These results also bear on a recent debate about whether corrections might actually make false information more familiar, thereby increasing the incidence of subsequent false beliefs (i.e., the familiarity backfire effect; Berinsky, 2017; Nyhan & Reifler, 2010; Schwarz et al., 2007; Skurnik, Yoon, Park, & Schwarz, 2005). In contrast to the backfire account, the latest research in this domain has indicated that explicit warnings or corrections of false statements actually have a small positive (and certainly not negative) impact on subsequent perceptions of accuracy (Ecker, Hogan, & Lewandowsky, 2017; Lewandowsky, Ecker, Seifert, Schwarz, & Cook, 2012; Pennycook & Rand, 2018b; Swire, Ecker, & Lewandowsky, 2017). In our data, the positive effect of a single prior exposure ($d = .20$ in Study 2) was effectively equivalent to the negative effect of the *disputed* warning ($d = .17$ in Study 2). Thus, although any benefit arising from the disputed tag is immediately wiped out by the prior exposure effect, we also did not find any evidence of a meaningful backfire. Our findings therefore support recent skepticism about the robustness and importance of the familiarity backfire effect.

Societal Implications

Our findings have important implications for the functioning of democracy, which relies on an informed electorate. Specif-

ically, our results shed some light on what can be done to combat belief in fake news. We employed a warning that was developed by Facebook to curb the influence of fake news on its social media platform (*Disputed by 3rd Party Fact-Checkers*). We found that this warning did not disrupt the illusory truth effect, an observation that resonates with findings of previous work demonstrating that, for example, explicitly labeling consumer claims as false (Skurnik et al., 2005) or retracting pieces of misinformation in news articles (Berinsky, 2017; Ecker et al., 2010; Lewandowsky et al., 2012; Nyhan & Reifler, 2010) are not necessarily effective strategies for decreasing long-term misperceptions (but see Swire, Ecker, & Lewandowsky, 2017). Nonetheless, it is important to note that the warning did successfully decrease subsequent overall perceptions of the accuracy of fake-news headlines; the warning's effect was just not specific to the particular fake-news headlines that the warning was attached to (and so the illusory truth effect survived the warning). Thus, the warning appears to have increased general skepticism, which increased the overall sensitivity to fake news (i.e., the warning decreased perceptions of fake-news accuracy without affecting judgments for real news). The warning also successfully decreased people's willingness to share fake-news headlines on social media. However, neither of these warning effect sizes were particularly large—for example, as described earlier, the negative impact of the warning on accuracy was entirely canceled out by the positive impact of repetition. That result, coupled with the persistence of the illusory truth effect we observed and the possibility of an “implied truth” effect whereby tagging some fake headlines may increase the perceived accuracy of untagged fake headlines (Pennycook & Rand, 2017), suggests that larger solutions are needed to prevent people from ever seeing fake news in the first place, rather than showing qualifiers aimed at making people discount the fake news that they do see.

Finally, our findings have implications beyond just fake news on social media. They suggest that politicians who continually repeat false statements will be successful, at least to some extent, in convincing people that those statements are in fact true. Indeed, the word *delusion* derives from a Latin term conveying the notion of mocking, defrauding, and deceiving. That the illusory truth effect is evident for highly salient and impactful information suggests that repetition may also play an important role in domains beyond politics, such as the formation of religious and paranormal beliefs where claims are difficult to either validate or reject empirically. When the truth is hard to come by, fluency is an attractive stand-in.

Context

In this research program, we used cognitive psychological theory and techniques to illuminate issues that have clear consequences for everyday life, with the hope of generating insights that are both practically and theoretically relevant. The topic of fake news—and disinformation more broadly—is of great relevance to current public discourse and policy making and fits squarely in the domain of cognitive psychology. Plainly, this topic is something that cognitive psychologists should be able to say something specific and illuminating about.

References

- Allcott, H., & Gentzkow, M. (2017). *Social media and fake news in the 2016 election* (NBER Working Paper No. 23098). Retrieved from <http://www.nber.org/papers/w23098>
- Alter, A. L., & Oppenheimer, D. M. (2009). Uniting the tribes of fluency to form a metacognitive nation. *Personality and Social Psychology Review*, 13, 219–235. <http://dx.doi.org/10.1177/1088868309341564>
- Arkes, H. R., Boehm, L. E., & Xu, G. (1991). Determinants of judged validity. *Journal of Experimental Social Psychology*, 27, 576–605. [http://dx.doi.org/10.1016/0022-1031\(91\)90026-3](http://dx.doi.org/10.1016/0022-1031(91)90026-3)
- Arkes, H. R., Hackett, C., & Boehm, L. (1989). The generality of the relation between familiarity and judged validity. *Journal of Behavioral Decision Making*, 2, 81–94. <http://dx.doi.org/10.1002/bdm.3960020203>
- Bacon, F. T. (1979). Credibility of repeated statements: Memory for trivia. *Journal of Experimental Psychology: Human Learning and Memory*, 5, 241–252. <http://dx.doi.org/10.1037/0278-7393.5.3.241>
- Beck, J. (2017, March 13). This article won't change your mind: The fact on why facts alone can't fight false beliefs. *Atlantic*. Retrieved from <https://www.theatlantic.com/science/archive/2017/03/this-article-wont-change-your-mind/519093/>
- Begg, I. M., Anas, A., & Farinacci, S. (1992). Dissociation of processes in belief: Source recollection, statement familiarity, and the illusion of truth. *Journal of Experimental Psychology: General*, 121, 446–458. <http://dx.doi.org/10.1037/0096-3445.121.4.446>
- Berinsky, A. A. J. (2017). Rumors and health care reform: Experiments in political misinformation. *British Journal of Political Science*, 47, 241–262. <http://dx.doi.org/10.1017/S0007123415000186>
- Calvert, D. (2017, March 6). *The psychology behind fake news: Cognitive biases help explain our polarized media climate*. Retrieved from <https://insight.kellogg.northwestern.edu/article/the-psychology-behind-fake-news>
- Coppock, A. (2016). *Generalizing from survey experiments conducted on Mechanical Turk: A replication approach*. Retrieved from https://alexandercoppock.files.wordpress.com/2016/02/coppock_generalizability2.pdf
- Corlett, P. R., Krystal, J. H., Taylor, J. R., & Fletcher, P. C. (2009). Why do delusions persist? *Frontiers in Human Neuroscience*, 3, 12. <http://dx.doi.org/10.3389/neuro.09.012.2009>
- Dechêne, A., Stahl, C., Hansen, J., & Wänke, M. (2010). The truth about the truth: A meta-analytic review of the truth effect. *Personality and Social Psychology Review*, 14, 238–257. <http://dx.doi.org/10.1177/1088868309352251>
- Diana, R., Yonelinas, A., & Ranganath, C. (2007). Imaging recollection and familiarity in the medial temporal lobe: A three-component model. *Trends in Cognitive Sciences*, 11, 379–386. <http://dx.doi.org/10.1016/j.tics.2007.08.001>
- Ecker, U. K. H., Hogan, J., & Lewandowsky, S. (2017). Reminders and repetition of misinformation: Helping or hindering its retraction? *Journal of Applied Research in Memory & Cognition*, 6, 185–192. <http://dx.doi.org/10.1016/j.jarmac.2017.01.014>
- Ecker, U. K. H., Lewandowsky, S., & Tang, D. T. W. (2010). Explicit warnings reduce but do not eliminate the continued influence of misinformation. *Memory & Cognition*, 38, 1087–1100. <http://dx.doi.org/10.3758/MC.38.8.1087>
- Fazio, L. K., Brashier, N. M., Payne, B. K., & Marsh, E. J. (2015). Knowledge does not protect against illusory truth. *Journal of Experimental Psychology: General*, 144, 993–1002. <http://dx.doi.org/10.1037/xge0000098>
- Flynn, D., Nyhan, B., & Reifler, J. (2017). The nature and origins of misperceptions: Understanding false and unsupported beliefs about politics. *Advances in Political Psychology*, 38(S1), 127–150. <http://dx.doi.org/10.1111/pops.12394>

- Gabielkov, M., Ramachandran, A., Chaintreau, A., & Legout, A. (2016). *Social clicks: What and who gets read on Twitter?* Retrieved from <http://dl.acm.org/citation.cfm?id=2901462>
- Goldman, R. (2016, December 24). Reading fake news, Pakistani minister directs nuclear threat at Israel. *The New York Times*. Retrieved from <https://www.nytimes.com/2016/12/24/world/asia/pakistan-israel-khawaja-asif-fake-news-nuclear.html>
- Gottfried, J., & Shearer, E. (2016). *News use across social media platforms 2016*. Retrieved from <http://www.journalism.org/2016/05/26/news-use-across-social-media-platforms-2016/>
- Hasher, L., Goldstein, D., & Toppino, T. (1977). Frequency and the conference of referential validity. *Journal of Verbal Learning and Verbal Behavior*, 16, 107–112. [http://dx.doi.org/10.1016/S0022-5371\(77\)80012-1](http://dx.doi.org/10.1016/S0022-5371(77)80012-1)
- Hawkins, S., & Hoch, S. (1992). Low-involvement learning: Memory without evaluation. *Journal of Consumer Research*, 19, 212–225. <http://dx.doi.org/10.1086/209297>
- Horton, J., Rand, D., & Zeckhauser, R. (2011). The online laboratory: Conducting experiments in a real labor market. *Experimental Economics*, 14, 399–425. <http://dx.doi.org/10.1007/s10683-011-9273-9>
- Johar, G., & Roggeveen, A. (2007). Changing false beliefs from repeated advertising: The role of claim-refutation alignment. *Journal of Consumer Psychology*, 17, 118–127. [http://dx.doi.org/10.1016/S1057-7408\(07\)70018-9](http://dx.doi.org/10.1016/S1057-7408(07)70018-9)
- Kahan, D. M. (2013). Ideology, motivated reasoning, and cognitive reflection. *Judgment and Decision Making*, 8, 407–424.
- Kahan, D. M. (2017, May 24). *Misconceptions, misinformation, and the logic of identity-protective cognition* (Cultural Cognition Project Working Paper Series No. 164; Yale Law School, Public Law Research Paper No. 605; Yale Law & Economics Research Paper No. 575). Available at <https://ssrn.com/abstract=2973067>
- Kahan, D. M., Peters, E., Wittlin, M., Slovic, P., Ouellette, L. L., Braman, D., & Mandel, G. (2012). The polarizing impact of science literacy and numeracy on perceived climate change risks. *Nature Climate Change*, 2, 732–735. <http://dx.doi.org/10.1038/nclimate1547>
- Krupnikov, Y., & Levine, A. (2014). Cross-sample comparisons and external validity. *Journal of Experimental Political Science*, 1, 59–80. <http://dx.doi.org/10.1017/xps.2014.7>
- Kunda, Z. (1990). The case for motivated reasoning. *Psychological Bulletin*, 108, 480–498. <http://dx.doi.org/10.1037/0033-2909.108.3.480>
- Lazer, D. M. J., Baum, M. A., Benkler, Y., Berinsky, A. J., Greenhill, K. M., Menczer, F., . . . Zittrain, J. L. (2018). The science of fake news. *Science*, 359, 1094–1096. <http://dx.doi.org/10.1126/science.aao2998>
- Lewandowsky, S., Ecker, U. K. H., Seifert, C. M., Schwarz, N., & Cook, J. (2012). Misinformation and its correction: Continued influence and successful debiasing. *Psychological Science in the Public Interest*, 13, 106–131. <http://dx.doi.org/10.1177/1529100612451018>
- Mercier, H., & Sperber, D. (2011). Why do humans reason? Arguments for an argumentative theory. *Behavioral and Brain Sciences*, 34, 57–74. <http://dx.doi.org/10.1017/S0140525X10000968>
- Mosseri, A. (2016, June 29). *Building a better news feed for you*. Retrieved from <http://newsroom.fb.com/news/2016/06/building-a-better-news-feed-for-you/>
- Mullinix, K., Leeper, T., Druckman, J., & Freese, J. (2015). The generalizability of survey experiments. *Journal of Experimental Political Science*, 2, 109–138. <http://dx.doi.org/10.1017/XPS.2015.19>
- Nyhan, B., & Reifler, J. (2010). When corrections fail: The persistence of political misperceptions. *Political Behavior*, 32, 303–330. <http://dx.doi.org/10.1007/s11109-010-9112-2>
- Pennycook, G., Fugelsang, J. A., & Koehler, D. J. (2015). What makes us think? A three-stage dual-process model of analytic engagement. *Cognitive Psychology*, 80, 34–72. <http://dx.doi.org/10.1016/j.cogpsych.2015.05.001>
- Pennycook, G., & Rand, D. (2017). *The implied truth effect: Attaching warnings to a subset of fake news stories increases perceived accuracy of stories without warnings*. Retrieved from https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3035384
- Pennycook, G., & Rand, D. G. (2018a). Lazy, not biased: Susceptibility to partisan fake news is better explained by lack of reasoning than by motivated reasoning. *Cognition*. <http://dx.doi.org/10.1016/j.cognition.2018.06.011>
- Pennycook, G., & Rand, D. G. (2018b). *Who falls for fake news? The roles of bullshit receptivity, overclaiming, familiarity, and analytic thinking*. Retrieved from https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3023545
- Polage, D. C. (2012). Making up history: False memories of fake news stories. *Europe's Journal of Psychology*, 8, 245–250. <http://dx.doi.org/10.5964/ejop.v8i2.456>
- Reber, R., Winkielman, P., & Schwarz, N. (1998). Effects of perceptual fluency on affective judgments. *Psychological Science*, 9, 45–48. <http://dx.doi.org/10.1111/1467-9280.00008>
- Redlawsk, D. (2002). Hot cognition or cool consideration? Testing the effects of motivated reasoning on political decision making. *Journal of Politics*, 64, 1021–1044. <http://dx.doi.org/10.1111/1468-2508.00161>
- Sanford, N., Veckenstedt, R., Moritz, S., Balzan, R. P., & Woodward, T. S. (2014). Impaired integration of disambiguating evidence in delusional schizophrenia patients. *Psychological Medicine*, 44, 2729–2738. <http://dx.doi.org/10.1017/S0033291714000397>
- Schwartz, M. (1982). Repetition and rated truth value of statements. *American Journal of Psychology*, 95, 393–407. <http://dx.doi.org/10.2307/1422132>
- Schwarz, N., Sanna, L. L. J., Skurnik, I., & Yoon, C. (2007). Metacognitive experiences and the intricacies of setting people straight: Implications for debiasing and public information campaigns. *Advances in Experimental Social Psychology*, 39, 127–161. [http://dx.doi.org/10.1016/S0065-2601\(06\)39003-X](http://dx.doi.org/10.1016/S0065-2601(06)39003-X)
- Shane, S. (2017, January 18). From headline to photograph, a fake news masterpiece. *New York Times*. Retrieved from <https://www.nytimes.com/2017/01/18/us/fake-news-hillary-clinton-cameron-harris.html>
- Silverman, C., Strapagiel, L., Shaban, H., & Hall, E. (2016, October 20). Hyperpartisan Facebook pages are publishing false and misleading information at an alarming rate. *Buzzfeed News*. Retrieved from <https://www.buzzfeed.com/craigsilverman/partisan-fb-pages-analysis>
- Singal, J. (2017, January 27). This is a great psychological framework for understanding how fake news spreads. *New York Magazine*. Retrieved from <http://nymag.com/scienceofus/2017/01/a-great-psychological-framework-for-understanding-fake-news.html>
- Skurnik, I., Yoon, C., Park, D. C., & Schwarz, N. (2005). How warnings about false claims become recommendations. *Journal of Consumer Research*, 31, 713–724. <http://dx.doi.org/10.1086/426605>
- Swire, B., Berinsky, A. J., Lewandowsky, S., & Ecker, U. K. H. (2017). Processing political misinformation: Comprehending the Trump phenomenon. *Royal Society Open Science*, 4, 160802. <http://dx.doi.org/10.1098/rsos.160802>
- Swire, B., Ecker, U. K. H., & Lewandowsky, S. (2017). The role of familiarity in correcting inaccurate information. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 43, 1948–1961. <http://dx.doi.org/10.1037/xlm0000422>
- Tauber, S., Dunlosky, J., & Rawson, K. (2013). General knowledge norms: Updated and expanded from the Nelson and Narens (1980) norms. *Behavior Research Methods*, 45, 1115–1143. <http://dx.doi.org/10.3758/s13428-012-0307-9>
- Unkelbach, C. (2007). Reversing the truth effect: Learning the interpretation of processing fluency in judgments of truth. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 33, 219–230. <http://dx.doi.org/10.1037/0278-7393.33.1.219>
- Unkelbach, C., & Rom, S. (2017). A referential theory of the repetition-induced truth effect. *Cognition*, 160, 110–126. <http://dx.doi.org/10.1016/j.cognition.2016.12.016>

Wang, W.-C., Brashier, N. M., Wing, E. A., Marsh, E. J., & Cabeza, R. (2016). On known unknowns: Fluency and the neural mechanisms of illusory truth. *Journal of Cognitive Neuroscience*, 28, 739–746. http://dx.doi.org/10.1162/jocn_a_00923

Watson, D., Clark, L. A., & Tellegen, A. (1988). Development and validation of brief measures of positive and negative affect: The PANAS scales. *Journal of Personality and Social Psychology*, 54, 1063–1070. <http://dx.doi.org/10.1037/0022-3514.54.6.1063>

Whittlesea, B. W. (1993). Illusions of familiarity. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 19, 1235–1253. <http://dx.doi.org/10.1037/0278-7393.19.6.1235>

Yonelinas, A. (2002). The nature of recollection and familiarity: A review of 30 years of research. *Journal of Memory and Language*, 46, 441–517. <http://dx.doi.org/10.1006/jmla.2002.2864>

Yonelinas, A., & Jacoby, L. (2012). The process-dissociation approach two decades later: Convergence, boundary conditions, and new directions. *Memory & Cognition*, 40, 663–680. <http://dx.doi.org/10.3758/s13421-012-0205-5>

Zajonc, R. (1968). Attitudinal effects of mere exposure. *Journal of Personality and Social Psychology*, 9, 1–27. <http://dx.doi.org/10.1037/h0025848>

Zajonc, R. (2001). Mere exposure: A gateway to the subliminal. *Current Directions in Psychological Science*, 10, 224–228. <http://dx.doi.org/10.1111/1467-8721.00154>

Appendix

Items Used in the Three Studies

Table A1

Study 1 Items

| Category | Type | Statement |
|----------|-------------------------------|---|
| Known | False (extremely implausible) | Smoking cigarettes is good for your lungs. The earth is a perfect square. Across the United States, only a total of 452 people voted in the last election. A single elephant weighs less than a single ant. |
| | True | More people live in the United States than in Malta. Cows are larger than sheep. Coffee is a more popular drink in America than goat milk. There are more than fifty stars in the universe. |
| Unknown | False | George was the name of the goldfish in the story of Pinocchio. Johnson was the last name of the man who killed Jesse James. Charles II was the first ruler of the Holy Roman Empire. Canopus is the name of the brightest star in the sky, excluding the sun. Tirpitz was the name of Germany's largest battleship that was sunk in World War II. John Kenneth Galbraith is the name of a well-known lawyer. Huxley is the name of the scientist who discovered radium. The Cotton Bowl takes place in Auston, Texas. The drachma is the monetary unit for Macedonia. |
| | True | Angel Falls is located in Brazil. The thigh bone is the largest bone in the human body. Bolivia borders the Pacific Ocean. The largest dam in the world is in Pakistan. Mexico is the world's largest producer of silver. More presidents of the United States were born in Virginia than any other state. Helsinki is the capital of Finland. Marconi is name of the inventor of the wireless radio. Billy the Kid's last name was Bonney. Tiber is the name of the river that runs through Rome. Canberra is the capital of Australia. |

Note. The unknown items were taken from Arkes, Hackett, & Boehm (1989); however, two of the “true” items (about Bolivia and Pakistan) are actually false. We retain the original labeling as it has no material effect on our results.

(Appendix continues)

Table A2
Study 2: Fake-News Items

| Political valence | Headline | Source |
|-------------------|---|--|
| Pro-Republican | Election Night: Hillary Was Drunk, Got Physical With Mook and Podesta Obama Was Going to Castro's Funeral—Until Trump Told Him This . . . Donald Trump Sent His Own Plane to Transport 200 Stranded Marines BLM Thug Protests President Trump With Selfie . . . Accidentally Shoots Himself in the Face NYT David Brooks: "Trump Needs to Decide if He Prefers to Resign, Be Impeached or Get Assassinated" Clint Eastwood Refuses to Accept Presidential Medal of Freedom From Obama, Says "He Is Not My President" | dailyheadlines.net thelastlineofdefense.org uconservative.com freedomdaily.com unitedstates-politics.com incredibleusnews.com |
| Pro-Democrat | Mike Pence: Gay Conversion Therapy Saved My Marriage Pennsylvania Federal Court Grants Legal Authority to Remove Trump After Russian Meddling Trump on Revamping the Military: We're Bringing Back the Draft FBI Director Comey Just Proved His Bias by Putting Trump Sign on His Front Lawn Sarah Palin Calls to Boycott Mall of America Because "Santa Was Always White in the Bible" Trump to Ban All TV Shows That Promote Gay Activity Starting With Empire as President | ncscooper.com bipartisanreport.com realnewsrighnow.com countercurrentnews.com politicono.com colossil.com |

Note. Fake- and real-news headlines as presented to participants can be found at the following link: <https://osf.io/txf46/>.

Table A3
Study 2: Real-News Items

| Political valence | Headline | Source |
|-------------------|---|---|
| Pro-Republican | Dems Scramble to Prevent Their Own From Defecting to Trump Majority of Americans Say Trump Can Keep Businesses, Poll Shows Donald Trump Strikes Conciliatory Tone in Meeting With Tech Executives Companies Are Already Canceling Plans to Move U.S. Jobs Abroad She Claimed She Was Attacked by Men Who Yelled "Trump" and Grabbed Her Hijab. Police Say She Lied. At GOP Convention Finale, Donald Trump Vows to Protect LGBTQ Community | foxnews.com bloomberg.com wsj.com msn.com washingtonpost.com fortune.com |
| Pro-Democrat | North Carolina Republicans Push Legislation to Hobble Incoming Democratic Governor Vladimir Putin "Personally Involved" in US Hack, Report Claims Trump Lashes Out at Vanity Fair, One Day After It Lambastes His Restaurant Donald Trump Says He'd "Absolutely" Require Muslims to Register The Small Businesses Near Trump Tower Are Experiencing a Miniature Recession Trump Questions Russia's Election Meddling on Twitter—Inaccurately | huffingtonpost.com theguardian.com npr.org nytimes.com slate.com nytimes.com |

Note. Fake- and real-news headlines as presented to participants can be found at the following link: <https://osf.io/txf46/>.

Table A4
Study 3: Fake-News Items

| Political valence | Headline | Source |
|-------------------|--|---|
| Pro-Republican | Election Night: Hillary Was Drunk, Got Physical With Mook and Podesta Donald Trump Protester Speaks Out: "I Was Paid \$3,500 to Protest Trump's Rally" NYT David Brooks: "Trump Needs to Decide if He Prefers to Resign, Be Impeached or Get Assassinated" Clint Eastwood Refuses to Accept Presidential Medal of Freedom From Obama, Says "He Is Not My President" | dailyheadlines.net abcnews.com.co unitedstates-politics.com incredibleusnews.com |
| Pro-Democrat | Donald Trump Sent His Own Plane to Transport 200 Stranded Marines BLM Thug Protests President Trump With Selfie . . . Accidentally Shoots Himself in the Face FBI Director Comey Just Proved His Bias by Putting Trump Sign on His Front Lawn Pennsylvania Federal Court Grants Legal Authority to Remove Trump After Russian Meddling Sarah Palin Calls to Boycott Mall of America Because "Santa Was Always White in the Bible" Trump to Ban All TV Shows That Promote Gay Activity Starting With Empire as President Mike Pence: Gay Conversion Therapy Saved My Marriage Trump on Revamping the Military: We're Bringing Back the Draft | uconservative.com freedomdaily.com countercurrentnews.com bipartisanreport.com politicono.com colossil.com ncscooper.com realnewsrighnow.com |

Note. Fake- and real-news headlines as presented to participants can be found at the following link: <https://osf.io/txf46/>.

(Appendix continues)

Table A5

Study 3: Real-News Items

| Political valence | Headline | Source |
|-------------------|--|--------------------|
| Pro-Republican | House Speaker Ryan Praises Trump for Maintaining Congressional Strength | cnn.com |
| | Donald Trump Strikes Conciliatory Tone in Meeting With Tech Executives | wsj.com |
| | At GOP Convention Finale, Donald Trump Vows to Protect LGBTQ Community | fortune.com |
| | Companies Are Already Canceling Plans to Move U.S. Jobs Abroad | msn.com |
| | Majority of Americans Say Trump Can Keep Businesses, Poll Shows | bloomberg.com |
| Pro-Democrat | She Claimed She Was Attacked by Men Who Yelled “Trump” and Grabbed Her Hijab. Police Say She Lied. | washingtonpost.com |
| | North Carolina Republicans Push Legislation to Hobble Incoming Democratic Governor | huffingtonpost.com |
| | Vladimir Putin “Personally Involved” in US Hack, Report Claims | theguardian.com |
| | Trump Lashes Out at Vanity Fair, One Day After It Lambastes His Restaurant | npr.org |
| | Trump Questions Russia’s Election Meddling on Twitter—Inaccurately | nytimes.com |
| | The Small Businesses Near Trump Tower Are Experiencing a Miniature Recession | slate.com |
| | Donald Trump Says He’d “Absolutely” Require Muslims to Register | nytimes.com |

Note. Fake- and real-news headlines as presented to participants can be found at the following link: <https://osf.io/txf46/>.

Received August 28, 2017

Revision received April 25, 2018

Accepted May 2, 2018 ■